

---

# Class Metaprogramming

Everyone knows that debugging is twice as hard as writing a program in the first place.  
So if you're as clever as you can be when you write it, how will you ever debug it?

—Brian W. Kernighan and P. J. Plauger, *The Elements of Programming Style*<sup>1</sup>

Class metaprogramming is the art of creating or customizing classes at runtime. Classes are first-class objects in Python, so a function can be used to create a new class at any time, without using the `class` keyword. Class decorators are also functions, but designed to inspect, change, and even replace the decorated class with another class. Finally, metaclasses are the most advanced tool for class metaprogramming: they let you create whole new categories of classes with special traits, such as the abstract base classes we've already seen.

Metaclasses are powerful, but hard to justify and even harder to get right. Class decorators solve many of the same problems and are easier to understand. Furthermore, Python 3.6 implemented [PEP 487—Simpler customization of class creation](#), providing special methods supporting tasks that previously required metaclasses or class decorators.<sup>2</sup>

This chapter presents the class metaprogramming techniques in ascending order of complexity.

---

<sup>1</sup> Quote from Chapter 2, “Expression” of *The Elements of Programming Style*, 2nd ed. (McGraw-Hill), page 10.

<sup>2</sup> That doesn't mean PEP 487 broke code that used those features. It just means that some code that used class decorators or metaclasses prior to Python 3.6 can now be refactored to use plain classes, resulting in simpler and possibly more efficient code.



This is an exciting topic, and it's easy to get carried away. So I must offer this advice.

For the sake of readability and maintainability, you should probably avoid the techniques described in this chapter in application code.

On the other hand, these are the tools of the trade if you want to write the next great Python framework.

## What's New in This Chapter

All the code in the “Class Metaprogramming” chapter of the first edition of *Fluent Python* still runs correctly. However, some of the previous examples no longer represent the simplest solutions in light of new features added since Python 3.6.

I replaced those examples with different ones, highlighting Python's new metaprogramming features or adding further requirements to justify the use of the more advanced techniques. Some of the new examples leverage type hints to provide class builders similar to the `@dataclass` decorator and `typing.NamedTuple`.

“[Metaclasses in the Real World](#)” on [page 947](#) is a new section with some high-level considerations about the applicability of metaclasses.



Some of the best refactorings are removing code made redundant by newer and simpler ways of solving the same problems. This applies to production code as well as books.

We'll get started by reviewing attributes and methods defined in the Python Data Model for all classes.

## Classes as Objects

Like most program entities in Python, classes are also objects. Every class has a number of attributes defined in the Python Data Model, documented in “[4.13. Special Attributes](#)” of the “Built-in Types” chapter in *The Python Standard Library*. Three of those attributes appeared several times in this book already: `__class__`, `__name__`, and `__mro__`. Other class standard attributes are:

`cls.__bases__`

The tuple of base classes of the class.

`cls.__qualname__`

The qualified name of a class or function, which is a dotted path from the global scope of the module to the class definition. This is relevant when the class is defined inside another class. For example, in a Django model class such as `Ox`, there is an inner class called `Meta`. The `__qualname__` of `Meta` is `Ox.Meta`, but its `__name__` is just `Meta`. The specification for this attribute is [PEP 3155—Qualified name for classes and functions](#).

`cls.__subclasses__()`

This method returns a list of the immediate subclasses of the class. The implementation uses weak references to avoid circular references between the superclass and its subclasses—which hold a strong reference to the superclasses in their `__bases__` attribute. The method lists subclasses currently in memory. Subclasses in modules not yet imported will not appear in the result.

`cls.mro()`

The interpreter calls this method when building a class to obtain the tuple of superclasses stored in the `__mro__` attribute of the class. A metaclass can override this method to customize the method resolution order of the class under construction.



None of the attributes mentioned in this section are listed by the `dir(...)` function.

Now, if a class is an object, what is the class of a class?

## type: The Built-In Class Factory

We usually think of `type` as a function that returns the class of an object, because that's what `type(my_object)` does: it returns `my_object.__class__`.

However, `type` is a class that creates a new class when invoked with three arguments.

Consider this simple class:

```
class MyClass(MySuperClass, MyMixin):
    x = 42

    def x2(self):
        return self.x * 2
```

Using the `type` constructor, you can create `MyClass` at runtime with this code:

```

MyClass = type('MyClass',
               (MySuperClass, MyMixin),
               {'x': 42, 'x2': lambda self: self.x * 2},
               )

```

That `type` call is functionally equivalent to the previous `class MyClass...` block statement.

When Python reads a class statement, it calls `type` to build the class object with these parameters:

**name**

The identifier that appears after the `class` keyword, e.g., `MyClass`.

**bases**

The tuple of superclasses given in parentheses after the class identifier, or `(object,)` if superclasses are not mentioned in the class statement.

**dict**

A mapping of attribute names to values. Callables become methods, as we saw in [“Methods Are Descriptors” on page 898](#). Other values become class attributes.



The `type` constructor accepts optional keyword arguments, which are ignored by `type` itself, but are passed untouched into `__init_subclass__`, which must consume them. We’ll study that special method in [“Introducing `\_\_init\_subclass\_\_`” on page 914](#), but I won’t cover the use of keyword arguments. For more, please read [PEP 487—Simpler customization of class creation](#).

The `type` class is a *metaclass*: a class that builds classes. In other words, instances of the `type` class are classes. The standard library provides a few other metaclasses, but `type` is the default:

```

>>> type(7)
<class 'int'>
>>> type(int)
<class 'type'>
>>> type(ZeroError)
<class 'type'>
>>> class Whatever:
...     pass
...
>>> type(Whatever)
<class 'type'>

```

We’ll build custom metaclasses in [“Metaclasses 101” on page 931](#).

Next, we’ll use the `type` built-in to make a function that builds classes.

# A Class Factory Function

The standard library has a class factory function that appears several times in this book: `collections.namedtuple`. In [Chapter 5](#) we also saw `typing.NamedTuple` and `@dataclass`. All of these class builders leverage techniques covered in this chapter.

We'll start with a super simple factory for classes of mutable objects—the simplest possible replacement for `@dataclass`.

Suppose I'm writing a pet shop application and I want to store data for dogs as simple records. But I don't want to write boilerplate like this:

```
class Dog:
    def __init__(self, name, weight, owner):
        self.name = name
        self.weight = weight
        self.owner = owner
```

Boring...each field name appears three times, and that boilerplate doesn't even buy us a nice repr:

```
>>> rex = Dog('Rex', 30, 'Bob')
>>> rex
<__main__.Dog object at 0x2865bac>
```

Taking a hint from `collections.namedtuple`, let's create a `record_factory` that creates simple classes like `Dog` on the fly. [Example 24-1](#) shows how it should work.

*Example 24-1. Testing `record_factory`, a simple class factory*

```
>>> Dog = record_factory('Dog', 'name weight owner') ❶
>>> rex = Dog('Rex', 30, 'Bob')
>>> rex ❷
Dog(name='Rex', weight=30, owner='Bob')
>>> name, weight, _ = rex ❸
>>> name, weight
('Rex', 30)
>>> "{2}'s dog weighs {1}kg".format(*rex) ❹
"Bob's dog weighs 30kg"
>>> rex.weight = 32 ❺
>>> rex
Dog(name='Rex', weight=32, owner='Bob')
>>> Dog.__mro__ ❻
(<class 'factories.Dog'>, <class 'object'>)
```

- ❶ Factory can be called like `namedtuple`: class name, followed by attribute names separated by spaces in a single string.
- ❷ Nice repr.

- ③ Instances are iterable, so they can be conveniently unpacked on assignment...
- ④ ...or when passing to functions like `format`.
- ⑤ A record instance is mutable.
- ⑥ The newly created class inherits from `object`—no relationship to our factory.

The code for `record_factory` is in [Example 24-2](#).<sup>3</sup>

*Example 24-2. `record_factory.py`: a simple class factory*

```

from typing import Union, Any
from collections.abc import Iterable, Iterator

FieldNames = Union[str, Iterable[str]] ①

def record_factory(cls_name: str, field_names: FieldNames) -> type[tuple]: ②

    slots = parse_identifiers(field_names) ③

    def __init__(self, *args, **kwargs) -> None: ④
        attrs = dict(zip(self.__slots__, args))
        attrs.update(kwargs)
        for name, value in attrs.items():
            setattr(self, name, value)

    def __iter__(self) -> Iterator[Any]: ⑤
        for name in self.__slots__:
            yield getattr(self, name)

    def __repr__(self): ⑥
        values = ', '.join(f'{name}={value!r}'
                           for name, value in zip(self.__slots__, self))
        cls_name = self.__class__.__name__
        return f'{cls_name}({values})'

    cls_attrs = dict( ⑦
        __slots__=slots,
        __init__=__init__,
        __iter__=__iter__,
        __repr__=__repr__,
    )

    return type(cls_name, (object,), cls_attrs) ⑧

```

---

<sup>3</sup> Thanks to my friend J. S. O. Bueno for contributing to this example.

```
def parse_identifiers(names: FieldNames) -> tuple[str, ...]:
    if isinstance(names, str):
        names = names.replace(',', ' ').split() ❸
    if not all(s.isidentifier() for s in names):
        raise ValueError('names must all be valid identifiers')
    return tuple(names)
```

- ❶ User can provide field names as a single string or an iterable of strings.
- ❷ Accept arguments like the first two of `collections.namedtuple`; return a type—i.e., a class that behaves like a `tuple`.
- ❸ Build a tuple of attribute names; this will be the `__slots__` attribute of the new class.
- ❹ This function will become the `__init__` method in the new class. It accepts positional and/or keyword arguments.<sup>4</sup>
- ❺ Yield the field values in the order given by `__slots__`.
- ❻ Produce the nice repr, iterating over `__slots__` and `self`.
- ❼ Assemble a dictionary of class attributes.
- ❽ Build and return the new class, calling the type constructor.
- ❾ Convert names separated by spaces or commas to list of `str`.

**Example 24-2** is the first time we’ve seen `type` in a type hint. If the annotation was just `-> type`, that would mean that `record_factory` returns a class—and it would be correct. But the annotation `-> type[tuple]` is more precise: it says the returned class will be a subclass of `tuple`.

The last line of `record_factory` in **Example 24-2** builds a class named by the value of `cls_name`, with `object` as its single immediate base class, and with a namespace loaded with `__slots__`, `__init__`, `__iter__`, and `__repr__`, of which the last three are instance methods.

We could have named the `__slots__` class attribute anything else, but then we’d have to implement `__setattr__` to validate the names of attributes being assigned,

---

<sup>4</sup> I did not add type hints to the arguments because the actual types are `Any`. I put the return type hint because otherwise `Mypy` will not check inside the method.

because for our record-like classes we want the set of attributes to be always the same and in the same order. However, recall that the main feature of `__slots__` is saving memory when you are dealing with millions of instances, and using `__slots__` has some drawbacks, discussed in “Saving Memory with `__slots__`” on page 384.



Instances of classes created by `record_factory` are not serializable—that is, they can’t be exported with the `dump` function from the `pickle` module. Solving this problem is beyond the scope of this example, which aims to show the `type` class in action in a simple use case. For the full solution, study the source code for `collections.namedtuple`; search for the word “pickling.”

Now let’s see how to emulate more modern class builders like `typing.NamedTuple`, which takes a user-defined class written as a `class` statement, and automatically enhances it with more functionality.

## Introducing `__init_subclass__`

Both `__init_subclass__` and `__set_name__` were proposed in [PEP 487—Simpler customization of class creation](#). We saw the `__set_name__` special method for descriptors for the first time in “[LineItem Take #4: Automatic Naming of Storage Attributes](#)” on page 887. Now let’s study `__init_subclass__`.

In [Chapter 5](#), we saw that `typing.NamedTuple` and `@dataclass` let programmers use the `class` statement to specify attributes for a new class, which is then enhanced by the class builder with the automatic addition of essential methods like `__init__`, `__repr__`, `__eq__`, etc.

Both of these class builders read type hints in the user’s `class` statement to enhance the class. Those type hints also allow static type checkers to validate code that sets or gets those attributes. However, `NamedTuple` and `@dataclass` do not take advantage of the type hints for attribute validation at runtime. The `Checked` class in the next example does.



It is not possible to support every conceivable static type hint for runtime type checking, which is probably why `typing.NamedTuple` and `@dataclass` don’t even try it. However, some types that are also concrete classes can be used with `Checked`. This includes simple types often used for field contents, such as `str`, `int`, `float`, and `bool`, as well as lists of those types.

[Example 24-3](#) shows how to use `Checked` to build a `Movie` class.



*Example 24-3. `initsub/checkedlib.py`: doctest for creating a `Movie` subclass of `Checked`*

```
>>> class Movie(Checked): ❶
...     title: str ❷
...     year: int
...     box_office: float
...
>>> movie = Movie(title='The Godfather', year=1972, box_office=137) ❸
>>> movie.title
'The Godfather'
>>> movie ❹
Movie(title='The Godfather', year=1972, box_office=137.0)
```

- ❶ `Movie` inherits from `Checked`—which we’ll define later in [Example 24-5](#).
- ❷ Each attribute is annotated with a constructor. Here I used built-in types.
- ❸ `Movie` instances must be created using keyword arguments.
- ❹ In return, you get a nice `__repr__`.

The constructors used as the attribute type hints may be any callable that takes zero or one argument and returns a value suitable for the intended field type, or rejects the argument by raising `TypeError` or `ValueError`.

Using built-in types for the annotations in [Example 24-3](#) means the values must be acceptable by the constructor of the type. For `int`, this means any `x` such that `int(x)` returns an `int`. For `str`, anything goes at runtime, because `str(x)` works with any `x` in Python.<sup>5</sup>

When called with no arguments, the constructor should return a default value of its type.<sup>6</sup>

This is standard behavior for Python’s built-in constructors:

```
>>> int(), float(), bool(), str(), list(), dict(), set()
(0, 0.0, False, '', [], {}, set())
```

---

<sup>5</sup> That’s true for any object, except when its class overrides the `__str__` or `__repr__` methods inherited from object with broken implementations.

<sup>6</sup> This solution avoids using `None` as a default. Avoiding null values is a [good idea](#). They are hard to avoid in general, but easy in some cases. In Python as well as SQL, I prefer to represent missing data in a text field with an empty string instead of `None` or `NULL`. Learning Go reinforced this idea: variables and struct fields of primitive types in Go are initialized by default with a “zero value.” See “[Zero values](#)” in the online *Tour of Go* if you are curious.

In a Checked subclass like `Movie`, missing parameters create instances with default values returned by the field constructors. For example:

```
>>> Movie(title='Life of Brian')
Movie(title='Life of Brian', year=0, box_office=0.0)
```

The constructors are used for validation during instantiation and when an attribute is set directly on an instance:

```
>>> blockbuster = Movie(title='Avatar', year=2009, box_office='billions')
Traceback (most recent call last):
...
TypeError: 'billions' is not compatible with box_office:float
>>> movie.year = 'MCMLXXII'
Traceback (most recent call last):
...
TypeError: 'MCMLXXII' is not compatible with year:int
```



### Checked Subclasses and Static Type Checking

In a `.py` source file with a `movie` instance of `Movie`, as defined in [Example 24-3](#), Mypy flags this assignment as a type error:

```
movie.year = 'MCMLXXII'
```

However, Mypy can't detect type errors in this constructor call:

```
blockbuster = Movie(title='Avatar', year='MMIX')
```

That's because `Movie` inherits `Checked.__init__`, and the signature of that method must accept any keyword arguments to support arbitrary user-defined classes.

On the other hand, if you declare a Checked subclass field with the type hint `list[float]`, Mypy can flag assignments of lists with incompatible contents, but `Checked` will ignore the type parameter and treat that the same as `list`.

Now let's look at the implementation of `checkedlib.py`. The first class is the `Field` descriptor, as shown in [Example 24-4](#).

*Example 24-4. `initsub/checkedlib.py`: the `Field` descriptor class*

```
from collections.abc import Callable ❶
from typing import Any, NoReturn, get_type_hints

class Field:
    def __init__(self, name: str, constructor: Callable) -> None: ❷
        if not callable(constructor) or constructor is type(None): ❸
            raise TypeError(f'{name!r} type hint must be callable')
        self.name = name
```

```

self.constructor = constructor

def __set__(self, instance: Any, value: Any) -> None:
    if value is ...: ❷
        value = self.constructor()
    else:
        try:
            value = self.constructor(value) ❸
        except (TypeError, ValueError) as e: ❹
            type_name = self.constructor.__name__
            msg = f'{value!r} is not compatible with {self.name}:{type_name}'
            raise TypeError(msg) from e
        instance.__dict__[self.name] = value ❺

```

- ❶ Recall that since Python 3.9, the Callable type for annotations is the ABC in `collections.abc`, and not the deprecated `typing.Callable`.
- ❷ This is a minimal Callable type hint; the parameter type and return type for `constructor` are both implicitly `Any`.
- ❸ For runtime checking, we use the `callable` built-in.<sup>7</sup> The test against `type(None)` is necessary because Python reads `None` in a type as `NoneType`, the class of `None` (therefore callable), but a useless constructor that only returns `None`.
- ❹ If `Checked.__init__` sets the value as `...` (the `Ellipsis` built-in object), we call the constructor with no arguments.
- ❺ Otherwise, call the constructor with the given value.
- ❻ If constructor raises either of these exceptions, we raise `TypeError` with a helpful message including the names of the field and constructor; e.g., `'MMIX' is not compatible with year:int`.
- ❼ If no exceptions were raised, the value is stored in the `instance.__dict__`.

In `__set__`, we need to catch `TypeError` and `ValueError` because built-in constructors may raise either of them, depending on the argument. For example, `float(None)` raises `TypeError`, but `float('A')` raises `ValueError`. On the other hand, `float('8')` raises no error and returns `8.0`. I hereby declare that this is a feature and not a bug of this toy example.

---

<sup>7</sup> I believe that `callable` should be made suitable for type hinting. As of May 6, 2021, this is an [open issue](#).



In “LineItem Take #4: Automatic Naming of Storage Attributes” on page 887, we saw the handy `__set_name__` special method for descriptors. We don’t need it in the `Field` class because the descriptors are not instantiated in client source code; the user declares types that are constructors, as we saw in the `Movie` class (Example 24-3). Instead, the `Field` descriptor instances are created at runtime by the `Checked.__init_subclass__` method, which we’ll see in Example 24-5.

Now let’s focus on the `Checked` class. I split it in two listings. Example 24-5 shows the top of the class, which includes the most important methods in this example. The remaining methods are in Example 24-6.

*Example 24-5. `initsub/checkedlib.py`: the most important methods of the `Checked` class*

```
class Checked:
    @classmethod
    def _fields(cls) -> dict[str, type]: ❶
        return get_type_hints(cls)

    def __init_subclass__(subclass) -> None: ❷
        super().__init_subclass__() ❸
        for name, constructor in subclass._fields().items(): ❹
            setattr(subclass, name, Field(name, constructor)) ❺

    def __init__(self, **kwargs: Any) -> None:
        for name in self._fields(): ❻
            value = kwargs.pop(name, ...) ❼
            setattr(self, name, value) ❽
        if kwargs: ❾
            self.__flag_unknown_attrs(*kwargs) ❿
```

- ❶ I wrote this class method to hide the call to `typing.get_type_hints` from the rest of the class. If I need to support Python  $\geq 3.10$  only, I’d call `inspect.get_annotations` instead. Review “Problems with Annotations at Runtime” on page 538 for the issues with those functions.
- ❷ `__init_subclass__` is called when a subclass of the current class is defined. It gets that new subclass as its first argument—which is why I named the argument `subclass` instead of the usual `cls`. For more on this, see “`__init_subclass__` Is Not a Typical Class Method” on page 919.
- ❸ `super().__init_subclass__()` is not strictly necessary, but should be invoked to play nice with other classes that might implement `__init_subclass__()` in

the same inheritance graph. See “Multiple Inheritance and Method Resolution Order” on page 494.

- ④ Iterate over each field name and constructor...
- ⑤ ...creating an attribute on subclass with that name bound to a Field descriptor parameterized with name and constructor.
- ⑥ For each name in the class fields...
- ⑦ ...get the corresponding value from kwargs and remove it from kwargs. Using ... (the Ellipsis object) as default allows us to distinguish between arguments given the value None from arguments that were not given.<sup>8</sup>
- ⑧ This setattr call triggers Checked.\_\_setattr\_\_, shown in Example 24-6.
- ⑨ If there are remaining items in kwargs, their names do not match any of the declared fields, and \_\_init\_\_ will fail.
- ⑩ The error is reported by \_\_flag\_unknown\_attrs, listed in Example 24-6. It takes a \*names argument with the unknown attribute names. I used a single asterisk in \*kwargs to pass its keys as a sequence of arguments.

### **`__init_subclass__` Is Not a Typical Class Method**

The @classmethod decorator is never used with `__init_subclass__`, but that doesn't mean much, because the `__new__` special method behaves as a class method even without @classmethod. The first argument that Python passes to `__init_subclass__` is a class. However, it is never the class where `__init_subclass__` is implemented: it is a newly defined subclass of that class. That's unlike `__new__` and every other class method that I know about. Therefore, I think `__init_subclass__` is not a class method in the usual sense, and it is misleading to name the first argument `cls`. The [\\_\\_init\\_subclass\\_\\_ documentation](#) names the argument `cls` but explains: “...called whenever the containing class is subclassed. `cls` is then the new subclass.”

Now let's see the remaining methods of the Checked class, continuing from Example 24-5. Note that I prepended `_` to the `_fields` and `_asdict` method names

---

<sup>8</sup> As mentioned in “Loops, Sentinels, and Poison Pills” on page 721, the Ellipsis object is a convenient and safe sentinel value. It has been around for a long time, but recently people are finding more uses for it, as we see in type hints and NumPy.

for the same reason the `collections.namedtuple` API does: to reduce the chance of name clashes with user-defined field names.

*Example 24-6. `initsub/checkedlib.py`: remaining methods of the `Checked` class*

```
def __setattr__(self, name: str, value: Any) -> None: ❶
    if name in self.__fields(): ❷
        cls = self.__class__
        descriptor = getattr(cls, name)
        descriptor.__set__(self, value) ❸
    else: ❹
        self.__flag_unknown_attrs(name)

def __flag_unknown_attrs(self, *names: str) -> NoReturn: ❺
    plural = 's' if len(names) > 1 else ''
    extra = ', '.join(f'{name!r}' for name in names)
    cls_name = repr(self.__class__.__name__)
    raise AttributeError(f'{cls_name} object has no attribute{plural} {extra}')

def __asdict(self) -> dict[str, Any]: ❻
    return {
        name: getattr(self, name)
        for name, attr in self.__class__.__dict__.items()
        if isinstance(attr, Field)
    }

def __repr__(self) -> str: ❼
    kwargs = ', '.join(
        f'{key}={value!r}' for key, value in self.__asdict().items()
    )
    return f'{self.__class__.__name__}({kwargs})'
```

- ❶ Intercept all attempts to set an instance attribute. This is needed to prevent setting an unknown attribute.
- ❷ If the attribute name is known, fetch the corresponding descriptor.
- ❸ Usually we don't need to call the descriptor `__set__` explicitly. It was necessary in this case because `__setattr__` intercepts all attempts to set an attribute on the instance, including in the presence of an overriding descriptor such as `Field`.<sup>9</sup>
- ❹ Otherwise, the attribute name is unknown, and an exception will be raised by `__flag_unknown_attrs`.

---

<sup>9</sup> The subtle concept of an overriding descriptor was explained in “[Overriding Descriptors](#)” on page 894.

- ⑤ Build a helpful error message listing all unexpected arguments, and raise `AttributeError`. This is a rare example of the `NoReturn` special type, covered in “[NoReturn](#)” on page 294.
- ⑥ Create a dict from the attributes of a `Movie` object. I’d call this method `_as_dict`, but I followed the convention started by the `_asdict` method in `collections.namedtuple`.
- ⑦ Implementing a nice `__repr__` is the main reason for having `_asdict` in this example.

The `Checked` example illustrates how to handle overriding descriptors when implementing `__setattr__` to block arbitrary attribute setting after instantiation. It is debatable whether implementing `__setattr__` is worthwhile in this example. Without it, setting `movie.director = 'Greta Gerwig'` would succeed, but the `director` attribute would not be checked in any way, and would not appear in the `__repr__` nor would it be included in the dict returned by `_asdict`—both defined in [Example 24-6](#).

In `record_factory.py` ([Example 24-2](#)) I solved this issue using the `__slots__` class attribute. However, this simpler solution is not viable in this case, as explained next.

## Why `__init_subclass__` Cannot Configure `__slots__`

The `__slots__` attribute is only effective if it is one of the entries in the class namespace passed to `type.__new__`. Adding `__slots__` to an existing class has no effect. Python invokes `__init_subclass__` only after the class is built—by then it’s too late to configure `__slots__`. A class decorator can’t configure `__slots__` either, because it is applied even later than `__init_subclass__`. We’ll explore these timing issues in “[What Happens When: Import Time Versus Runtime](#)” on page 925.

To configure `__slots__` at runtime, your own code must build the class namespace passed as the last argument of `type.__new__`. To do that, you can write a class factory function, like `record_factory.py`, or you can take the nuclear option and implement a metaclass. We will see how to dynamically configure `__slots__` in “[Metaclasses 101](#)” on page 931.

Before [PEP 487](#) simplified the customization of class creation with `__init_subclass__` in Python 3.7, similar functionality had to be implemented using a class decorator. That’s the focus of the next section.

# Enhancing Classes with a Class Decorator

A class decorator is a callable that behaves similarly to a function decorator: it gets the decorated class as an argument, and should return a class to replace the decorated class. Class decorators often return the decorated class itself, after injecting more methods in it via attribute assignment.

Probably the most common reason to choose a class decorator over the simpler `__init_subclass__` is to avoid interfering with other class features, such as inheritance and metaclasses.<sup>10</sup>

In this section, we'll study *checkeddeco.py*, which provides the same service as *checkedlib.py*, but using a class decorator. As usual, we'll start by looking at a usage example, extracted from the doctests in *checkeddeco.py* (Example 24-7).

*Example 24-7. checkeddeco.py: creating a Movie class decorated with @checked*

```
>>> @checked
... class Movie:
...     title: str
...     year: int
...     box_office: float
...
>>> movie = Movie(title='The Godfather', year=1972, box_office=137)
>>> movie.title
'The Godfather'
>>> movie
Movie(title='The Godfather', year=1972, box_office=137.0)
```

The only difference between Example 24-7 and Example 24-3 is the way the `Movie` class is declared: it is decorated with `@checked` instead of subclassing `Checked`. Otherwise, the external behavior is the same, including the type validation and default value assignments shown after Example 24-3 in “Introducing `__init_subclass__`” on page 914.

Now let's look at the implementation of *checkeddeco.py*. The imports and `Field` class are the same as in *checkedlib.py*, listed in Example 24-4. There is no other class, only functions in *checkeddeco.py*.

The logic previously implemented in `__init_subclass__` is now part of the `checked` function—the class decorator listed in Example 24-8.

---

<sup>10</sup> This rationale appears in the abstract of [PEP 557–Data Classes](#) to explain why it was implemented as a class decorator.



Example 24-8. *checkeddeco.py*: the class decorator

```
def checked(cls: type) -> type: ❶
    for name, constructor in _fields(cls).items(): ❷
        setattr(cls, name, Field(name, constructor)) ❸

    cls._fields = classmethod(_fields) # type: ignore ❹

    instance_methods = ( ❺
        __init__,
        __repr__,
        __setattr__,
        __asdict__,
        __flag_unknown_attrs,
    )
    for method in instance_methods: ❻
        setattr(cls, method.__name__, method)

    return cls ❼
```

- ❶ Recall that classes are instances of `type`. These type hints strongly suggest this is a class decorator: it takes a class and returns a class.
- ❷ `_fields` is a top-level function defined later in the module (in [Example 24-9](#)).
- ❸ Replacing each attribute returned by `_fields` with a `Field` descriptor instance is what `__init_subclass__` did in [Example 24-5](#). Here there is more work to do...
- ❹ Build a class method from `_fields`, and add it to the decorated class. The `type: ignore` comment is needed because Mypy complains that `type` has no `_fields` attribute.
- ❺ Module-level functions that will become instance methods of the decorated class.
- ❻ Add each of the `instance_methods` to `cls`.
- ❼ Return the decorated `cls`, fulfilling the essential contract of a class decorator.

Every top-level function in *checkeddeco.py* is prefixed with an underscore, except the checked decorator. This naming convention makes sense for a couple of reasons:

- checked is part of the public interface of the *checkeddeco.py* module, but the other functions are not.
- The functions in [Example 24-9](#) will be injected in the decorated class, and the leading `_` reduces the chance of naming conflicts with user-defined attributes and methods of the decorated class.

The rest of *checkeddeco.py* is listed in [Example 24-9](#). Those module-level functions have the same code as the corresponding methods of the Checked class of *checked-lib.py*. They were explained in [Examples 24-5](#) and [24-6](#).

Note that the `_fields` function does double duty in *checkeddeco.py*. It is used as a regular function in the first line of the checked decorator, and it will also be injected as a class method of the decorated class.

*Example 24-9. checkeddeco.py: the methods to be injected in the decorated class*

```
def _fields(cls: type) -> dict[str, type]:
    return get_type_hints(cls)

def __init__(self: Any, **kwargs: Any) -> None:
    for name in self._fields():
        value = kwargs.pop(name, ...)
        setattr(self, name, value)
    if kwargs:
        self.__flag_unknown_attrs(*kwargs)

def __setattr__(self: Any, name: str, value: Any) -> None:
    if name in self._fields():
        cls = self.__class__
        descriptor = getattr(cls, name)
        descriptor.__set__(self, value)
    else:
        self.__flag_unknown_attrs(name)

def __flag_unknown_attrs(self: Any, *names: str) -> NoReturn:
    plural = 's' if len(names) > 1 else ''
    extra = ', '.join(f'{name!r}' for name in names)
    cls_name = repr(self.__class__.__name__)
    raise AttributeError(f'{cls_name} has no attribute{plural} {extra}')

def __asdict(self: Any) -> dict[str, Any]:
    return {
        name: getattr(self, name)
        for name, attr in self.__class__.__dict__.items()
        if isinstance(attr, Field)
```

```

    }

def __repr__(self: Any) -> str:
    kwargs = ', '.join(
        f'{key}={value!r}' for key, value in self._asdict().items()
    )
    return f'{self.__class__.__name__}({kwargs})'

```

The *checkeddeco.py* module implements a simple but usable class decorator. Python’s `@dataclass` does a lot more. It supports many configuration options, adds more methods to the decorated class, handles or warns about conflicts with user-defined methods in the decorated class, and even traverses the `__mro__` to collect user-defined attributes declared in the superclasses of the decorated class. The [source code](#) of the `dataclasses` package in Python 3.9 is more than 1,200 lines long.

For metaprogramming classes, we must be aware of when the Python interpreter evaluates each block of code during the construction of a class. This is covered next.

## What Happens When: Import Time Versus Runtime

Python programmers talk about “import time” versus “runtime,” but the terms are not strictly defined and there is a gray area between them.

At import time, the interpreter:

1. Parses the source code of a *.py* module in one pass from top to bottom. This is when a `SyntaxError` may occur.
2. Compiles the bytecode to be executed.
3. Executes the top-level code of the compiled module.

If there is an up-to-date *.pyc* file available in the local `__pycache__`, parsing and compiling are skipped because the bytecode is ready to run.

Although parsing and compiling are definitely “import time” activities, other things may happen at that time, because almost every statement in Python is executable in the sense that they can potentially run user code and may change the state of the user program.

In particular, the `import` statement is not merely a declaration,<sup>11</sup> but it actually runs all the top-level code of a module when it is imported for the first time in the process. Further imports of the same module will use a cache, and then the only effect will be binding the imported objects to names in the client module. That top-level code may

---

11 Contrast with the `import` statement in Java, which is just a declaration to let the compiler know that certain packages are required.

do anything, including actions typical of “runtime,” such as writing to a log or connecting to a database.<sup>12</sup> That’s why the border between “import time” and “runtime” is fuzzy: the `import` statement can trigger all sorts of “runtime” behavior. Conversely, “import time” can also happen deep inside runtime, because the `import` statement and the `__import__()` built-in can be used inside any regular function.

This is all rather abstract and subtle, so let’s do some experiments to see what happens when.

## Evaluation Time Experiments

Consider an *evaldemo.py* script that uses a class decorator, a descriptor, and a class builder based on `__init_subclass__`, all defined in a *builderlib.py* module. The modules have several `print` calls to show what happens under the covers. Otherwise, they don’t perform anything useful. The goal of these experiments is to observe the order in which these `print` calls happen.



Applying a class decorator and a class builder with `__init_subclass__` together in single class is likely a sign of overengineering or desperation. This unusual combination is useful in these experiments to show the timing of the changes that a class decorator and `__init_subclass__` can apply to a class.

Let’s start by checking out *builderlib.py*, split into two parts: [Example 24-10](#) and [Example 24-11](#).

*Example 24-10. builderlib.py: top of the module*

```
print('@ builderlib module start')

class Builder: ❶
    print('@ Builder body')

    def __init_subclass__(cls): ❷
        print(f'@ Builder.__init_subclass__({cls!r})')

        def inner_0(self): ❸
            print(f'@ SuperA.__init_subclass__:inner_0({self!r})')

        cls.method_a = inner_0

    def __init__(self):
```

---

<sup>12</sup> I’m not saying opening a database connection just because a module is imported is a good idea, only pointing out it can be done.

```

        super().__init__()
        print(f'@ Builder.__init__({self!r})')

def deco(cls): ❷
    print(f'@ deco({cls!r})')

    def inner_1(self): ❸
        print(f'@ deco:inner_1({self!r})')

    cls.method_b = inner_1
    return cls ❹

```

- ❶ This is a class builder to implement...
- ❷ ...an `__init_subclass__` method.
- ❸ Define a function to be added to the subclass in the assignment below.
- ❹ A class decorator.
- ❺ Function to be added to the decorated class.
- ❻ Return the class received as an argument.

Continuing with *builderlib.py* in [Example 24-11](#)...

*Example 24-11. builderlib.py: bottom of the module*

```

class Descriptor: ❶
    print('@ Descriptor body')

    def __init__(self): ❷
        print(f'@ Descriptor.__init__({self!r})')

    def __set_name__(self, owner, name): ❸
        args = (self, owner, name)
        print(f'@ Descriptor.__set_name__({args!r})')

    def __set__(self, instance, value): ❹
        args = (self, instance, value)
        print(f'@ Descriptor.__set__({args!r})')

    def __repr__(self):
        return '<Descriptor instance>'

print('@ builderlib module end')

```

- ❶ A descriptor class to demonstrate when...
- ❷ ...a descriptor instance is created, and when...
- ❸ ...`__set_name__` will be invoked during the owner class construction.
- ❹ Like the other methods, this `__set__` doesn't do anything except display its arguments.

If you import *builderlib.py* in the Python console, this is what you get:

```
>>> import builderlib
@ builderlib module start
@ Builder body
@ Descriptor body
@ builderlib module end
```

Note that the lines printed by *builderlib.py* are prefixed with `@`.

Now let's turn to *evaldemo.py*, which will trigger special methods in *builderlib.py* (Example 24-12).

*Example 24-12. evaldemo.py: script to experiment with builderlib.py*

```
#!/usr/bin/env python3

from builderlib import Builder, deco, Descriptor

print('# evaldemo module start')

@deco ❶
class Klass(Builder): ❷
    print('# Klass body')

    attr = Descriptor() ❸

    def __init__(self):
        super().__init__()
        print(f'# Klass.__init__({self!r})')

    def __repr__(self):
        return '<Klass instance>'

def main(): ❹
    obj = Klass()
    obj.method_a()
    obj.method_b()
    obj.attr = 999
```

```

if __name__ == '__main__':
    main()

print('# evaldemo module end')

```

- ❶ Apply a decorator.
- ❷ Subclass Builder to trigger its `__init_subclass__`.
- ❸ Instantiate the descriptor.
- ❹ This will only be called if the module is run as the main program.

The print calls in *evaldemo.py* show a # prefix. If you open the console again and import *evaldemo.py*, [Example 24-13](#) is the output.

*Example 24-13. Console experiment with evaldemo.py*

```

>>> import evaldemo
@ builderlib module start ❶
@ Builder body
@ Descriptor body
@ builderlib module end
# evaldemo module start
# Class body ❷
@ Descriptor.__init__(<Descriptor instance>) ❸
@ Descriptor.__set_name__(<Descriptor instance>,
    <class 'evaldemo.Klass'>, 'attr') ❹
@ Builder.__init_subclass__(<class 'evaldemo.Klass'>) ❺
@ deco(<class 'evaldemo.Klass'>) ❻
# evaldemo module end

```

- ❶ The top four lines are the result of `from builderlib import...`. They will not appear if you didn't close the console after the previous experiment, because *builderlib.py* is already loaded.
- ❷ This signals that Python started reading the body of `Klass`. At this point, the class object does not exist yet.
- ❸ The descriptor instance is created and bound to `attr` in the namespace that Python will pass to the default class object constructor: `type.__new__`.
- ❹ At this point, Python's built-in `type.__new__` has created the `Klass` object and calls `__set_name__` on each descriptor instance of descriptor classes that provide that method, passing `Klass` as the owner argument.

- ⑤ `type.__new__` then calls `__init_subclass__` on the superclass of `Klass`, passing `Klass` as the single argument.
- ⑥ When `type.__new__` returns the class object, Python applies the decorator. In this example, the class returned by `deco` is bound to `Klass` in the module namespace.

The implementation of `type.__new__` is written in C. The behavior I just described is documented in the “[Creating the class object](#)” section of Python’s “[Data Model](#)” reference.

Note that the `main()` function of `evaldemo.py` ([Example 24-12](#)) was not executed in the console session ([Example 24-13](#)), therefore no instance of `Klass` was created. All the action we saw was triggered by “import time” operations: importing `builderlib` and defining `Klass`.

If you run `evaldemo.py` as a script, you will see the same output as [Example 24-13](#) with extra lines right before the end. The extra lines are the result of running `main()` ([Example 24-14](#)).

#### *Example 24-14. Running `evaldemo.py` as a program*

```
$ ./evaldemo.py
[... 9 lines omitted ...]
@ deco(<class '__main__.Klass'>) ①
@ Builder.__init__(<Klass instance>) ②
# Klass.__init__(<Klass instance>)
@ SuperA.__init_subclass__:inner_0(<Klass instance>) ③
@ deco:inner_1(<Klass instance>) ④
@ Descriptor.__set__(<Descriptor instance>, <Klass instance>, 999) ⑤
# evaldemo module end
```

- ① The top 10 lines—including this one—are the same as shown in [Example 24-13](#).
- ② Triggered by `super().__init__()` in `Klass.__init__`.
- ③ Triggered by `obj.method_a()` in `main`; `method_a` was injected by `SuperA.__init_subclass__`.
- ④ Triggered by `obj.method_b()` in `main`; `method_b` was injected by `deco`.
- ⑤ Triggered by `obj.attr = 999` in `main`.

A base class with `__init_subclass__` and a class decorator are powerful tools, but they are limited to working with a class already built by `type.__new__` under the



covers. In the rare occasions when you need to adjust the arguments passed to `type.__new__`, you need a metaclass. That's the final destination of this chapter—and this book.

## Metaclasses 101

[Metaclasses] are deeper magic than 99% of users should ever worry about. If you wonder whether you need them, you don't (the people who actually need them know with certainty that they need them, and don't need an explanation about why).

—Tim Peters, inventor of the Timsort algorithm and prolific Python contributor<sup>13</sup>

A metaclass is a class factory. In contrast with `record_factory` from [Example 24-2](#), a metaclass is written as a class. In other words, a metaclass is a class whose instances are classes. [Figure 24-1](#) depicts a metaclass using the Mills & Gizmos Notation: a mill producing another mill.




Figure 24-1. A metaclass is a class that builds classes.

Consider the Python object model: classes are objects, therefore each class must be an instance of some other class. By default, Python classes are instances of `type`. In other words, `type` is the metaclass for most built-in and user-defined classes:

```
>>> str.__class__  
<class 'type'>  
>>> from bulkfood_v5 import LineItem  
>>> LineItem.__class__  
<class 'type'>
```

---

<sup>13</sup> Message to `comp.lang.python`, subject: “[Acrimony in c.l.p.](#)”. This is another part of the same message from December 23, 2002, quoted in the [Preface](#). The TimBot was inspired that day.

```
>>> type.__class__
<class 'type'>
```

To avoid infinite regress, the class of `type` is `type`, as the last line shows.

Note that I am not saying that `str` or `LineItem` are subclasses of `type`. What I am saying is that `str` and `LineItem` are instances of `type`. They all are subclasses of `object`. [Figure 24-2](#) may help you confront this strange reality.




Figure 24-2. Both diagrams are true. The left one emphasizes that `str`, `type`, and `LineItem` are subclasses of `object`. The right one makes it clear that `str`, `object`, and `LineItem` are instances of `type`, because they are all classes.



The classes `object` and `type` have a unique relationship: `object` is an instance of `type`, and `type` is a subclass of `object`. This relationship is “magic”: it cannot be expressed in Python because either class would have to exist before the other could be defined. The fact that `type` is an instance of itself is also magical.

The next snippet shows that the class of `collections.Iterable` is `abc.ABCMeta`. Note that `Iterable` is an abstract class, but `ABCMeta` is a concrete class—after all, `Iterable` is an instance of `ABCMeta`:

```
>>> from collections.abc import Iterable
>>> Iterable.__class__
<class 'abc.ABCMeta'>
>>> import abc
>>> from abc import ABCMeta
>>> ABCMeta.__class__
<class 'type'>
```

Ultimately, the class of `ABCMeta` is also `type`. Every class is an instance of `type`, directly or indirectly, but only metaclasses are also subclasses of `type`. That’s the most important relationship to understand metaclasses: a metaclass, such as `ABCMeta`,

inherits from `type` the power to construct classes. Figure 24-3 illustrates this crucial relationship.

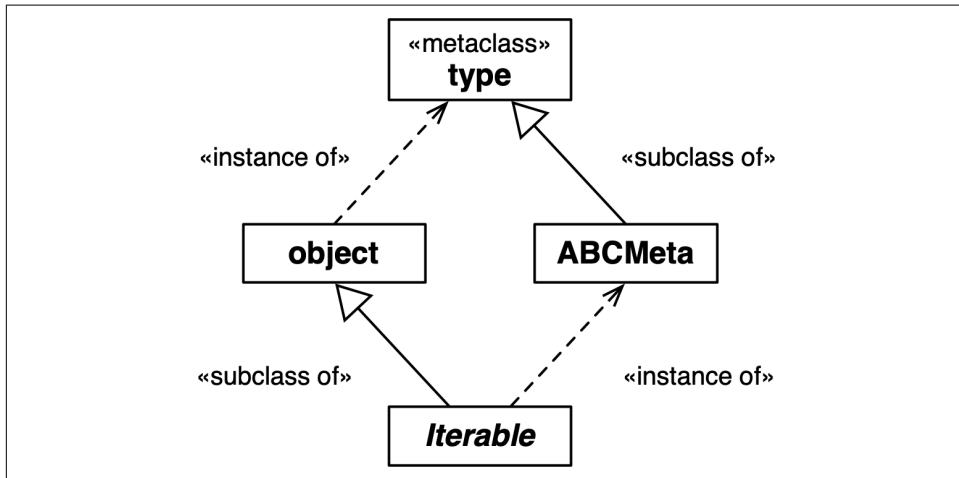


Figure 24-3. *Iterable* is a subclass of *object* and an instance of *ABCMeta*. Both *object* and *ABCMeta* are instances of *type*, but the key relationship here is that *ABCMeta* is also a subclass of *type*, because *ABCMeta* is a metaclass. In this diagram, *Iterable* is the only abstract class.

The important takeaway here is that metaclasses are subclasses of `type`, and that’s what makes them work as class factories. A metaclass can customize its instances by implementing special methods, as the next sections demonstrate.

## How a Metaclass Customizes a Class

To use a metaclass, it’s critical to understand how `__new__` works on any class. This was discussed in “Flexible Object Creation with `__new__`” on page 843.

The same mechanics happen at a “meta” level when a metaclass is about to create a new instance, which is a class. Consider this declaration:

```
class KClass(SuperKlass, metaclass=MetaKlass):
    x = 42
    def __init__(self, y):
        self.y = y
```

To process that `class` statement, Python calls `MetaKlass.__new__` with these arguments:

`meta_cls`

The metaclass itself (`MetaKlass`), because `__new__` works as class method.

`cls_name`

The string `Klass`.

`bases`

The single-element tuple `(SuperKlass,)`, with more elements in the case of multiple inheritance.

`cls_dict`

A mapping like:

```
{x: 42, '__init__': <function __init__ at 0x1009c4040>}
```

When you implement `MetaKlass.__new__`, you can inspect and change those arguments before passing them to `super().__new__`, which will eventually call `type.__new__` to create the new class object.

After `super().__new__` returns, you can also apply further processing to the newly created class before returning it to Python. Python then calls `SuperKlass.__init_subclass__`, passing the class you created, and then applies a class decorator to it, if one is present. Finally, Python binds the class object to its name in the surrounding namespace—usually the global namespace of a module, if the `class` statement was a top-level statement.

The most common processing made in a metaclass `__new__` is to add or replace items in the `cls_dict`—the mapping that represents the namespace of the class under construction. For instance, before calling `super().__new__`, you can inject methods in the class under construction by adding functions to `cls_dict`. However, note that adding methods can also be done after the class is built, which is why we were able to do it using `__init_subclass__` or a class decorator.

One attribute that you must add to the `cls_dict` before `type.__new__` runs is `__slots__`, as discussed in “[Why \\_\\_init\\_subclass\\_\\_ Cannot Configure \\_\\_slots\\_\\_](#)” on [page 921](#). The `__new__` method of a metaclass is the ideal place to configure `__slots__`. The next section shows how to do that.

## A Nice Metaclass Example

The `MetaBunch` metaclass presented here is a variation of the last example in Chapter 4 of *Python in a Nutshell, 3rd ed.*, by Alex Martelli, Anna Ravenscroft, and Steve

Holden, written to run on Python 2.7 and 3.5.<sup>14</sup> Assuming Python 3.6 or later, I was able to further simplify the code.

First, let's see what the Bunch base class provides:

```
>>> class Point(Bunch):
...     x = 0.0
...     y = 0.0
...     color = 'gray'
...
>>> Point(x=1.2, y=3, color='green')
Point(x=1.2, y=3, color='green')
>>> p = Point()
>>> p.x, p.y, p.color
(0.0, 0.0, 'gray')
>>> p
Point()
```

Remember that Checked assigns names to the Field descriptors in subclasses based on class variable type hints, which do not actually become attributes on the class since they don't have values.

Bunch subclasses, on the other hand, use actual class attributes with values, which then become the default values of the instance attributes. The generated `__repr__` omits the arguments for attributes that are equal to the defaults.

MetaBunch—the metaclass of Bunch—generates `__slots__` for the new class from the class attributes declared in the user's class. This blocks the instantiation and later assignment of undeclared attributes:

```
>>> Point(x=1, y=2, z=3)
Traceback (most recent call last):
...
AttributeError: No slots left for: 'z'
>>> p = Point(x=21)
>>> p.y = 42
>>> p
Point(x=21, y=42)
>>> p.flavor = 'banana'
Traceback (most recent call last):
...
AttributeError: 'Point' object has no attribute 'flavor'
```

---

<sup>14</sup> The authors kindly gave me permission to use their example. MetaBunch first appeared in a message posted by Martelli in the comp.lang.python group on July 7, 2002, with the subject line “a nice metaclass example (was Re: structs in python)”, following a discussion about record-like data structures in Python. Martelli's original code for Python 2.2 still runs after a single change: to use a metaclass in Python 3, you must use the metaclass keyword argument in the class declaration, e.g., `Bunch(metaclass=MetaBunch)`, instead of the older convention of adding a `__metaclass__` class-level attribute.

Now let's dive into the elegant code of MetaBunch in [Example 24-15](#).

*Example 24-15. metabunch/from3.6/bunch.py: MetaBunch metaclass and Bunch class*

```
class MetaBunch(type): ❶
    def __new__(meta_cls, cls_name, bases, cls_dict): ❷

        defaults = {} ❸

        def __init__(self, **kwargs): ❹
            for name, default in defaults.items(): ❺
                setattr(self, name, kwargs.pop(name, default))
            if kwargs: ❻
                extra = ', '.join(kwargs)
                raise AttributeError(f'No slots left for: {extra!r}')

        def __repr__(self): ❼
            rep = ', '.join(f'{name}={value!r}'
                            for name, default in defaults.items()
                            if (value := getattr(self, name)) != default)
            return f'{cls_name}({rep})'

        new_dict = dict(__slots__=[], __init__=__init__, __repr__=__repr__) ❽

        for name, value in cls_dict.items(): ❾
            if name.startswith('__') and name.endswith('__'): ❿
                if name in new_dict:
                    raise AttributeError(f"Can't set {name!r} in {cls_name!r}")
                new_dict[name] = value
            else: ⓫
                new_dict['__slots__'].append(name)
                defaults[name] = value
        return super().__new__(meta_cls, cls_name, bases, new_dict) ⓫

class Bunch(metaclass=MetaBunch): ⓫
    pass
```

- ❶ To create a new metaclass, inherit from type.
- ❷ `__new__` works as a class method, but the class is a metaclass, so I like to name the first argument `meta_cls` (`mcs` is a common alternative). The remaining three arguments are the same as the three-argument signature for calling `type()` directly to create a class.
- ❸ `defaults` will hold a mapping of attribute names and their default values.
- ❹ This will be injected into the new class.

- ⑤ Read the defaults and set the corresponding instance attribute with a value popped from kwargs or a default.
- ⑥ If there is still any item in kwargs, it means there are no slots left where we can place them. We believe in *failing fast* as best practice, so we don't want to silently ignore extra items. A quick and effective solution is to pop one item from kwargs and try to set it on the instance, triggering an `AttributeError` on purpose.
- ⑦ `__repr__` returns a string that looks like a constructor call—e.g., `Point(x=3)`, omitting the keyword arguments with default values.
- ⑧ Initialize namespace for the new class.
- ⑨ Iterate over the namespace of the user's class.
- ⑩ If a dunder name is found, copy the item to the new class namespace, unless it's already there. This prevents users from overwriting `__init__`, `__repr__`, and other attributes set by Python, such as `__qualname__` and `__module__`.
- ⑪ If not a dunder name, append to `__slots__` and save its value in defaults.
- ⑫ Build and return the new class.
- ⑬ Provide a base class, so users don't need to see `MetaBunch`.

`MetaBunch` works because it is able to configure `__slots__` before calling `super().__new__` to build the final class. As usual when metaprogramming, understanding the sequence of actions is key. Let's do another evaluation time experiment, now with a metaclass.

## Metaclass Evaluation Time Experiment

This is a variation of “[Evaluation Time Experiments](#)” on page 926, adding a metaclass to the mix. The `builderlib.py` module is the same as before, but the main script is now `evaldemo_meta.py`, listed in [Example 24-16](#).

*Example 24-16. evaldemo\_meta.py: experimenting with a metaclass*

```
#!/usr/bin/env python3

from builderlib import Builder, deco, Descriptor
from metalib import MetaClass ①

print('# evaldemo_meta module start')
```

```

@deco
class Klass(Builder, metaclass=MetaKlass): ❷
    print('# Klass body')

    attr = Descriptor()

    def __init__(self):
        super().__init__()
        print(f'# Klass.__init__({self!r})')

    def __repr__(self):
        return '<Klass instance>'

def main():
    obj = Klass()
    obj.method_a()
    obj.method_b()
    obj.method_c() ❸
    obj.attr = 999

if __name__ == '__main__':
    main()

print('# evaldemo_meta module end')

```

- ❶ Import MetaKlass from *metalib.py*, which we'll see in [Example 24-18](#).
- ❷ Declare Klass as a subclass of Builder and an instance of MetaKlass.
- ❸ This method is injected by MetaKlass.\_\_new\_\_, as we'll see.



In the interest of science, [Example 24-16](#) defies all reason and applies three different metaprogramming techniques together on Klass: a decorator, a base class using `__init_subclass__`, and a custom metaclass. If you do this in production code, please don't blame me. Again, the goal is to observe the order in which the three techniques interfere in the class construction process.

As in the previous evaluation time experiment, this example does nothing but print messages revealing the flow of execution. [Example 24-17](#) shows the code for the top part of *metalib.py*—the rest is in [Example 24-18](#).

*Example 24-17. metalib.py: the NosyDict class*

```

print('% metalib module start')

```



```
import collections

class NosyDict(collections.UserDict):
    def __setitem__(self, key, value):
        args = (self, key, value)
        print(f'% NosyDict.__setitem__ {args!r}')
        super().__setitem__(key, value)

    def __repr__(self):
        return '<NosyDict instance>'
```

I wrote the `NosyDict` class to override `__setitem__` to display each key and value as they are set. The metaclass will use a `NosyDict` instance to hold the namespace of the class under construction, revealing more of Python's inner workings.

The main attraction of *metatlib.py* is the metaclass in [Example 24-18](#). It implements the `__prepare__` special method, a class method that Python only invokes on metaclasses. The `__prepare__` method provides the earliest opportunity to influence the process of creating a new class.



When coding a metaclass, I find it useful to adopt this naming convention for special method arguments:

- Use `cls` instead of `self` for instance methods, because the instance is a class.
- Use `meta_cls` instead of `cls` for class methods, because the class is a metaclass. Recall that `__new__` behaves as a class method even without the `@classmethod` decorator.

*Example 24-18. metatlib.py: the MetaKlass*

```
class MetaKlass(type):
    print('% MetaKlass body')

    @classmethod ❶
    def __prepare__(meta_cls, cls_name, bases): ❷
        args = (meta_cls, cls_name, bases)
        print(f'% MetaKlass.__prepare__ {args!r}')
        return NosyDict() ❸

    def __new__(meta_cls, cls_name, bases, cls_dict): ❹
        args = (meta_cls, cls_name, bases, cls_dict)
        print(f'% MetaKlass.__new__ {args!r}')
        def inner_2(self):
            print(f'% MetaKlass.__new__:inner_2({self!r})')

        cls = super().__new__(meta_cls, cls_name, bases, cls_dict.data) ❺
```

```

        cls.method_c = inner_2 ❹

    return cls ❺

def __repr__(cls): ❸
    cls_name = cls.__name__
    return f"<class {cls_name!r} built by MetaKlass>"

print('% metalib module end')

```

- ❶ `__prepare__` should be declared as a class method. It is not an instance method because the class under construction does not exist yet when Python calls `__prepare__`.
- ❷ Python calls `__prepare__` on a metaclass to obtain a mapping to hold the namespace of the class under construction.
- ❸ Return `NosyDict` instance to be used as the namespace.
- ❹ `cls_dict` is a `NosyDict` instance returned by `__prepare__`.
- ❺ `type.__new__` requires a real dict as the last argument, so I give it the `data` attribute of `NosyDict`, inherited from `UserDict`.
- ❻ Inject a method in the newly created class.
- ❼ As usual, `__new__` must return the object just created—in this case, the new class.
- ❽ Defining `__repr__` on a metaclass allows customizing the `repr()` of class objects.

The main use case for `__prepare__` before Python 3.6 was to provide an `OrderedDict` to hold the attributes of the class under construction, so that the metaclass `__new__` could process those attributes in the order in which they appear in the source code of the user’s class definition. Now that dict preserves the insertion order, `__prepare__` is rarely needed. You will see a creative use for it in “A Metaclass Hack with `__prepare__`” on page 950.

Importing *metalib.py* in the Python console is not very exciting. Note the use of % to prefix the lines output by this module:

```

>>> import metalib
% metalib module start
% MetaKlass body
% metalib module end

```

Lots of things happen if you import *evaldemo\_meta.py*, as you can see in [Example 24-19](#).

*Example 24-19. Console experiment with evaldemo\_meta.py*

```
>>> import evaldemo_meta
@ builderlib module start
@ Builder body
@ Descriptor body
@ builderlib module end
% metalib module start
% MetaKlass body
% metalib module end
# evaldemo_meta module start ❶
% MetaKlass.__prepare__(<class 'metalib.MetaKlass'>, 'Klass', ❷
                        (<class 'builderlib.Builder'>,))
% NosyDict.__setitem__(<NosyDict instance>, '__module__', 'evaldemo_meta') ❸
% NosyDict.__setitem__(<NosyDict instance>, '__qualname__', 'Klass')
# Klass body
@ Descriptor.__init__(<Descriptor instance>) ❹
% NosyDict.__setitem__(<NosyDict instance>, 'attr', <Descriptor instance>) ❺
% NosyDict.__setitem__(<NosyDict instance>, '__init__',
                      <function Klass.__init__ at ...>) ❻
% NosyDict.__setitem__(<NosyDict instance>, '__repr__',
                      <function Klass.__repr__ at ...>)
% NosyDict.__setitem__(<NosyDict instance>, '__classcell__', <cell at ...: empty>)
% MetaKlass.__new__(<class 'metalib.MetaKlass'>, 'Klass',
                  (<class 'builderlib.Builder'>,), <NosyDict instance>) ❼
@ Descriptor.__set_name__(<Descriptor instance>,
                        <class 'Klass' built by MetaKlass>, 'attr') ❽
@ Builder.__init_subclass__(<class 'Klass' built by MetaKlass>)
@ deco(<class 'Klass' built by MetaKlass>)
# evaldemo_meta module end
```

- ❶ The lines before this are the result of importing *builderlib.py* and *metalib.py*.
- ❷ Python invokes `__prepare__` to start processing a class statement.
- ❸ Before parsing the class body, Python adds the `__module__` and `__qualname__` entries to the namespace of the class under construction.
- ❹ The descriptor instance is created...
- ❺ ...and bound to `attr` in the class namespace.
- ❻ `__init__` and `__repr__` methods are defined and added to the namespace.
- ❼ Once Python finishes processing the class body, it calls `MetaKlass.__new__`.
- ❽ `__set_name__`, `__init_subclass__`, and the decorator are invoked in this order, after the `__new__` method of the metaclass returns the newly constructed class.

If you run *evaldemo\_meta.py* as script, `main()` is called, and a few more things happen (Example 24-20).

*Example 24-20. Running evaldemo\_meta.py as a program*

```
$ ./evaldemo_meta.py
[... 20 lines omitted ...]
@ deco(<class 'Klass' built by MetaKlass>) ❶
@ Builder.__init__(<Klass instance>)
# Klass.__init__(<Klass instance>)
@ SuperA.__init_subclass__:inner_0(<Klass instance>)
@ deco:inner_1(<Klass instance>)
% MetaKlass.__new__:inner_2(<Klass instance>) ❷
@ Descriptor.__set__(<Descriptor instance>, <Klass instance>, 999)
# evaldemo_meta module end
```

- ❶ The top 21 lines—including this one—are the same shown in Example 24-19.
- ❷ Triggered by `obj.method_c()` in `main`; `method_c` was injected by `MetaKlass.__new__`.

Let's now go back to the idea of the `Checked` class with the `Field` descriptors implementing runtime type validation, and see how it can be done with a metaclass.

## A Metaclass Solution for Checked

I don't want to encourage premature optimization and overengineering, so here is a make-believe scenario to justify rewriting *checkedlib.py* with `__slots__`, which requires the application of a metaclass. Feel free to skip it.

### A Bit of Storytelling

Our *checkedlib.py* using `__init_subclass__` is a company-wide success, and our production servers have millions of instances of `Checked` subclasses in memory at any one time.

Profiling a proof-of-concept, we discover that using `__slots__` will reduce the cloud hosting bill for two reasons:

- Lower memory usage, as `Checked` instances don't need their own `__dict__`
- Higher performance, by removing `__setattr__`, which was created just to block unexpected attributes, but is triggered at instantiation and for all attribute setting before `Field.__set__` is called to do its job

The *metaclass/checkedlib.py* module we'll study next is a drop-in replacement for *init-sub/checkedlib.py*. The doctests embedded in them are identical, as well as the *checkedlib\_test.py* files for *pytest*.

The complexity in *checkedlib.py* is abstracted away from the user. Here is the source code of a script using the package:

```
from checkedlib import Checked

class Movie(Checked):
    title: str
    year: int
    box_office: float

if __name__ == '__main__':
    movie = Movie(title='The Godfather', year=1972, box_office=137)
    print(movie)
    print(movie.title)
```

That concise *Movie* class definition leverages three instances of the *Field* validating descriptor, a `__slots__` configuration, five methods inherited from *Checked*, and a metaclass to put it all together. The only visible part of *checkedlib* is the *Checked* base class.

Consider [Figure 24-4](#). The Mills & Gizmos Notation complements the UML class diagram by making the relationship between classes and instances more visible.

For example, a *Movie* class using the new *checkedlib.py* is an instance of *CheckedMeta*, and a subclass of *Checked*. Also, the *title*, *year*, and *box\_office* class attributes of *Movie* are three separate instances of *Field*. Each *Movie* instance has its own *\_title*, *\_year*, and *\_box\_office* attributes, to store the values of the corresponding fields.

Now let's study the code, starting with the *Field* class, shown in [Example 24-21](#).

The *Field* descriptor class is now a bit different. In the previous examples, each *Field* descriptor instance stored its value in the managed instance using an attribute of the same name. For example, in the *Movie* class, the *title* descriptor stored the field value in a *title* attribute in the managed instance. This made it unnecessary for *Field* to provide a `__get__` method.

However, when a class like *Movie* uses `__slots__`, it cannot have class attributes and instance attributes with the same name. Each descriptor instance is a class attribute, and now we need separate per-instance storage attributes. The code uses the descriptor name prefixed with a single `_`. Therefore *Field* instances have separate *name* and *storage\_name* attributes, and we implement *Field.\_\_get\_\_*.

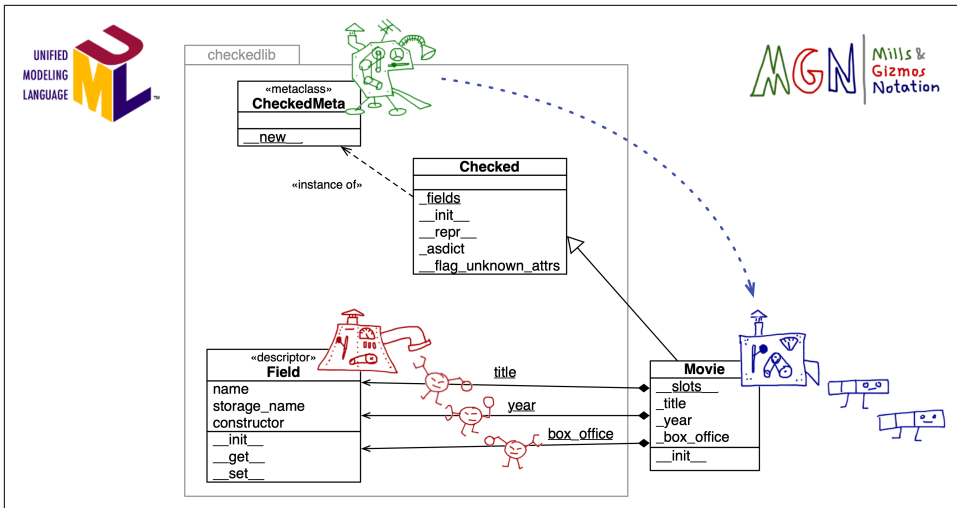


Figure 24-4. UML class diagram annotated with MGN: the *CheckedMeta* meta-mill builds the *Movie* mill. The *Field* mill builds the *title*, *year*, and *box\_office* descriptors, which are class attributes of *Movie*. The per-instance data for the fields is stored in the *\_title*, *\_year*, and *\_box\_office* instance attributes of *Movie*. Note the package boundary of *checkedlib*. The developer of *Movie* doesn't need to grok all the machinery inside *checkedlib.py*.

**Example 24-21** shows the source code for *Field*, with callouts describing only the changes in this version.

*Example 24-21. metaclass/checkedlib.py: the *Field* descriptor with *storage\_name* and *\_\_get\_\_**

```

class Field:
    def __init__(self, name: str, constructor: Callable) -> None:
        if not callable(constructor) or constructor is type(None):
            raise TypeError(f'{name!r} type hint must be callable')
        self.name = name
        self.storage_name = '_' + name ❶
        self.constructor = constructor

    def __get__(self, instance, owner=None):
        if instance is None: ❷
            return self
        return getattr(instance, self.storage_name) ❸

    def __set__(self, instance: Any, value: Any) -> None:
        if value is ...:
            value = self.constructor()
        else:

```

```

try:
    value = self.constructor(value)
except (TypeError, ValueError) as e:
    type_name = self.constructor.__name__
    msg = f'{value!r} is not compatible with {self.name}:{type_name}'
    raise TypeError(msg) from e
setattr(instance, self.storage_name, value) ❷

```

- ❶ Compute `storage_name` from the `name` argument.
- ❷ If `__get__` gets `None` as the instance argument, the descriptor is being read from the managed class itself, not a managed instance. So we return the descriptor.
- ❸ Otherwise, return the value stored in the attribute named `storage_name`.
- ❹ `__set__` now uses `setattr` to set or update the managed attribute.

Example 24-22 shows the code for the metaclass that drives this example.

*Example 24-22. metaclass/checkedlib.py: the `CheckedMeta` metaclass*

```

class CheckedMeta(type):

    def __new__(meta_cls, cls_name, bases, cls_dict): ❶
        if '__slots__' not in cls_dict: ❷
            slots = []
            type_hints = cls_dict.get('__annotations__', {}) ❸
            for name, constructor in type_hints.items(): ❹
                field = Field(name, constructor) ❺
                cls_dict[name] = field ❻
                slots.append(field.storage_name) ❼
            cls_dict['__slots__'] = slots ❽

        return super().__new__(
            meta_cls, cls_name, bases, cls_dict) ❾

```

- ❶ `__new__` is the only method implemented in `CheckedMeta`.
- ❷ Only enhance the class if its `cls_dict` doesn't include `__slots__`. If `__slots__` is already present, assume it is the `Checked` base class and not a user-defined subclass, and build the class as is.
- ❸ To get the type hints in prior examples, we used `typing.get_type_hints`, but that requires an existing class as the first argument. At this point, the class we are configuring does not exist yet, so we need to retrieve the `__annotations__`

directly from the `cls_dict`—the namespace of the class under construction, which Python passes as the last argument to the metaclass `__new__`.

- ④ Iterate over `type_hints` to...
- ⑤ ...build a `Field` for each annotated attribute...
- ⑥ ...overwrite the corresponding entry in `cls_dict` with the `Field` instance...
- ⑦ ...and append the `storage_name` of the field in the list we'll use to...
- ⑧ ...populate the `__slots__` entry in `cls_dict`—the namespace of the class under construction.
- ⑨ Finally, we call `super().__new__`.

The last part of *metaclass/checkedlib.py* is the `Checked` base class that users of this library will subclass to enhance their classes, like `Movie`.

The code for this version of `Checked` is the same as `Checked` in *initsub/checkedlib.py* (listed in [Example 24-5](#) and [Example 24-6](#)), with three changes:

1. Added an empty `__slots__` to signal to `CheckedMeta.__new__` that this class doesn't require special processing.
2. Removed `__init_subclass__`. Its job is now done by `CheckedMeta.__new__`.
3. Removed `__setattr__`. It became redundant because adding `__slots__` to the user-defined class prevents setting undeclared attributes.

[Example 24-23](#) is a complete listing of the final version of `Checked`.

*Example 24-23. metaclass/checkedlib.py: the `Checked` base class*

```
class Checked(metaclass=CheckedMeta):
    __slots__ = () # skip CheckedMeta.__new__ processing

    @classmethod
    def _fields(cls) -> dict[str, type]:
        return get_type_hints(cls)

    def __init__(self, **kwargs: Any) -> None:
        for name in self._fields():
            value = kwargs.pop(name, ...)
            setattr(self, name, value)
        if kwargs:
            self.__flag_unknown_attrs(*kwargs)
```



```

def __flag_unknown_attrs(self, *names: str) -> NoReturn:
    plural = 's' if len(names) > 1 else ''
    extra = ', '.join(f'{name!r}' for name in names)
    cls_name = repr(self.__class__.__name__)
    raise AttributeError(f'{cls_name} object has no attribute{plural} {extra}')

def __asdict(self) -> dict[str, Any]:
    return {
        name: getattr(self, name)
        for name, attr in self.__class__.__dict__.items()
        if isinstance(attr, Field)
    }

def __repr__(self) -> str:
    kwargs = ', '.join(
        f'{key}={value!r}' for key, value in self.__asdict().items()
    )
    return f'{self.__class__.__name__}({kwargs})'

```

This concludes the third rendering of a class builder with validated descriptors.

The next section covers some general issues related to metaclasses.

## Metaclasses in the Real World

Metaclasses are powerful, but tricky. Before deciding to implement a metaclass, consider the following points.

### Modern Features Simplify or Replace Metaclasses

Over time, several common use cases of metaclasses were made redundant by new language features:

#### *Class decorators*

Simpler to understand than metaclasses, and less likely to cause conflicts with base classes and metaclasses.

#### `__set_name__`

Avoids the need for custom metaclass logic to automatically set the name of a descriptor.<sup>15</sup>

---

<sup>15</sup> In the first edition of *Fluent Python*, the more advanced versions of the `LineItem` class used a metaclass just to set the storage name of the attributes. See the code in the metaclasses of [bulkfood in the first edition code repository](#).

`__init_subclass__`

Provides a way to customize class creation that is transparent to the end user and even simpler than a decorator—but may introduce conflicts in a complex class hierarchy.

*Built-in dict preserving key insertion order*

Eliminated the #1 reason to use `__prepare__`: to provide an `OrderedDict` to store the namespace of the class under construction. Python only calls `__prepare__` on metaclasses, so if you needed to process the class namespace in the order it appears in the source code, you had to use a metaclass before Python 3.6.

As of 2021, every actively maintained version of CPython supports all the features just listed.

I keep advocating these features because I see too much unnecessary complexity in our profession, and metaclasses are a gateway to complexity.

## Metaclasses Are Stable Language Features

Metaclasses were introduced in Python 2.2 in 2002, together with so-called “new-style classes,” descriptors, and properties.

It is remarkable that the `MetaBunch` example, first posted by Alex Martelli in July 2002, still works in Python 3.9—the only change being the way to specify the metaclass to use, which in Python 3 is done with the syntax `class Bunch(metaclass=MetaBunch):`.

None of the additions I mentioned in “[Modern Features Simplify or Replace Metaclasses](#)” on page 947 broke existing code using metaclasses. But legacy code using metaclasses can often be simplified by leveraging those features, especially if you can drop support to Python versions before 3.6—which are no longer maintained.

## A Class Can Only Have One Metaclass

If your class declaration involves two or more metaclasses, you will see this puzzling error message:

```
TypeError: metaclass conflict: the metaclass of a derived class
must be a (non-strict) subclass of the metaclasses of all its bases
```

This may happen even without multiple inheritance. For example, a declaration like this could trigger that `TypeError`:

```
class Record(abc.ABC, metaclass=PersistentMeta):
    pass
```

We saw that `abc.ABC` is an instance of the `abc.ABCMeta` metaclass. If that `Persistent` metaclass is not itself a subclass of `abc.ABCMeta`, you get a metaclass conflict.

There are two ways of dealing with that error:

- Find some other way of doing what you need to do, while avoiding at least one of the metaclasses involved.
- Write your own `PersistentABCMeta` metaclass as a subclass of both `abc.ABCMeta` and `PersistentMeta`, using multiple inheritance, and use that as the only metaclass for `Record`.<sup>16</sup>



I can imagine the solution of the metaclass with two base metaclasses implemented to meet a deadline. In my experience, metaclass programming always takes longer than anticipated, which makes this approach risky before a hard deadline. If you do it and make the deadline, the code may contain subtle bugs. Even in the absence of known bugs, you should consider this approach as technical debt simply because it is hard to understand and maintain.

## Metaclasses Should Be Implementation Details

Besides `type`, there are only six metaclasses in the entire Python 3.9 standard library. The better known metaclasses are probably `abc.ABCMeta`, `typing.NamedTupleMeta`, and `enum.EnumMeta`. None of them are intended to appear explicitly in user code. We may consider them implementation details.

Although you can do some really wacky metaprogramming with metaclasses, it's best to heed the **principle of least astonishment** so that most users can indeed regard metaclasses as implementation details.<sup>17</sup>

In recent years, some metaclasses in the Python standard library were replaced by other mechanisms, without breaking the public API of their packages. The simplest way to future-proof such APIs is to offer a regular class that users subclass to access the functionality provided by the metaclass, as we've done in our examples.

To wrap up our coverage of class metaprogramming, I will share with you the coolest, small example of metaclass I found as I researched this chapter.

---

<sup>16</sup> If you just got dizzy considering the implications of multiple inheritance with metaclasses, good for you. I'd stay way from this solution as well.

<sup>17</sup> I made a living writing Django code for a few years before I decided to study how Django's model fields were implemented. Only then I learned about descriptors and metaclasses.

## A Metaclass Hack with `__prepare__`

When I updated this chapter for the second edition, I needed to find simple but illuminating examples to replace the *bulkfood* `LineItem` code that no longer require metaclasses since Python 3.6.

The simplest and most interesting metaclass idea was given to me by João S. O. Bueno—better known as JS in the Brazilian Python community. One application of his idea is to create a class that autogenerates numeric constants:

```
>>> class Flavor(AutoConst):
...     banana
...     coconut
...     vanilla
...
>>> Flavor.vanilla
2
>>> Flavor.banana, Flavor.coconut
(0, 1)
```

Yes, that code works as shown! That’s actually a doctest in *autoconst\_demo.py*.

Here is the user-friendly `AutoConst` base class and the metaclass behind it, implemented in *autoconst.py*:

```
class AutoConstMeta(type):
    def __prepare__(name, bases, **kwargs):
        return WilyDict()

class AutoConst(metaclass=AutoConstMeta):
    pass
```

That’s it.

Clearly the trick is in `WilyDict`.

When Python processes the namespace of the user’s class and reads `banana`, it looks up that name in the mapping provided by `__prepare__`: an instance of `WilyDict`. `WilyDict` implements `__missing__`, covered in “The `__missing__` Method” on page 91. The `WilyDict` instance initially has no `'banana'` key, so the `__missing__` method is triggered. It makes an item on the fly with the key `'banana'` and the value `0`, returning that value. Python is happy with that, then tries to retrieve `'coconut'`. `WilyDict` promptly adds that entry with the value `1`, returning it. The same happens with `'vanilla'`, which is then mapped to `2`.

We've seen `__prepare__` and `__missing__` before. The real innovation is how JS put them together.

Here is the source code for `WilyDict`, also from *autoconst.py*:

```
class WilyDict(dict):
    def __init__(self, *args, **kwargs):
        super().__init__(*args, **kwargs)
        self.__next_value = 0

    def __missing__(self, key):
        if key.startswith('__') and key.endswith('__'):
            raise KeyError(key)
        self[key] = value = self.__next_value
        self.__next_value += 1
        return value
```

While experimenting, I found that Python looked up `__name__` in the namespace of the class under construction, causing `WilyDict` to add a `__name__` entry, and increment `__next_value`. So I added that if statement in `__missing__` to raise `KeyError` for keys that look like dunder attributes.

The *autoconst.py* package both requires and illustrates mastery of Python's dynamic class building machinery.

I had a great time adding more functionality to `AutoConstMeta` and `AutoConst`, but instead of sharing my experiments, I will let you have fun playing with JS's ingenious hack.

Here are some ideas:

- Make it possible to retrieve the constant name if you have the value. For example, `Flavor[2]` could return `'vanilla'`. You can do this by implementing `__getitem__` in `AutoConstMeta`. Since Python 3.9, you can implement `__class_getitem__` in `AutoConst` itself.
- Support iteration over the class, by implementing `__iter__` on the metaclass. It would make the `__iter__` yield the constants as `(name, value)` pairs.
- Implement a new `Enum` variant. This would be a major undertaking, because the `enum` package is full of tricks, including the `EnumMeta` metaclass with hundreds of lines of code and a nontrivial `__prepare__` method.

Enjoy!



The `__class_getitem__` special method was added in Python 3.9 to support generic types, as part of [PEP 585—Type Hinting Generics In Standard Collections](#). Thanks to `__class_getitem__`, Python’s core developers did not have to write a new metaclass for the built-in types to implement `__getitem__` so that we could write generic type hints like `list[int]`. This is a narrow feature, but representative of a wider use case for metaclasses: implementing operators and other special methods to work at the class level, such as making the class itself iterable, just like Enum subclasses.

## Wrapping Up

Metaclasses, as well as class decorators and `__init_subclass__` are useful for:

- Subclass registration
- Subclass structural validation
- Applying decorators to many methods at once
- Object serialization
- Object-relational mapping
- Object-based persistence
- Implementing special methods at the class level
- Implementing class features found in other languages, such as [traits](#) and [aspect-oriented programming](#)

Class metaprogramming can also help with performance issues in some cases, by performing tasks at import time that otherwise would execute repeatedly at runtime.

To wrap up, let’s recall Alex Martelli’s final advice from his essay [“Waterfowl and ABCs” on page 443](#):

And, *don’t* define custom ABCs (or metaclasses) in production code... if you feel the urge to do so, I’d bet it’s likely to be a case of “all problems look like a nail”-syndrome for somebody who just got a shiny new hammer—you (and future maintainers of your code) will be much happier sticking with straightforward and simple code, eschewing such depths.

I believe Martelli’s advice applies not only to ABCs and metaclasses, but also to class hierarchies, operator overloading, function decorators, descriptors, class decorators, and class builders using `__init_subclass__`.

Those powerful tools exist primarily to support library and framework development. Applications naturally should *use* those tools, as provided by the Python standard

library or external packages. But *implementing* them in application code is often premature abstraction.

Good frameworks are extracted, not invented.<sup>18</sup>

—David Heinemeier Hansson, creator of Ruby on Rails

## Chapter Summary

This chapter started with an overview of attributes found in class objects, such as `__qualname__` and the `__subclasses__()` method. Next, we saw how the `type` builtin can be used to construct classes at runtime.

The `__init_subclass__` special method was introduced, with the first iteration of a Checked base class designed to replace attribute type hints in user-defined subclasses with `Field` instances that apply constructors to enforce the type of those attributes at runtime.

The same idea was implemented with a `@checked` class decorator that adds features to user-defined classes, similar to what `__init_subclass__` allows. We saw that neither `__init_subclass__` nor a class decorator can dynamically configure `__slots__`, because they operate only after a class is created.

The concepts of “import time” and “runtime” were clarified with experiments showing the order in which Python code is executed when modules, descriptors, class decorators, and `__init_subclass__` is involved.

Our coverage of metaclasses began with an overall explanation of `type` as a metaclass, and how user-defined metaclasses can implement `__new__` to customize the classes it builds. We then saw our first custom metaclass, the classic `MetaBunch` example using `__slots__`. Next, another evaluation time experiment demonstrated how the `__prepare__` and `__new__` methods of a metaclass are invoked earlier than `__init_subclass__` and class decorators, providing opportunities for deeper class customization.

The third iteration of a Checked class builder with `Field` descriptors and custom `__slots__` configuration was presented, followed by some general considerations about metaclass usage in practice.

Finally, we saw the `AutoConst` hack invented by João S. O. Bueno, based on the cunning idea of a metaclass with `__prepare__` returning a mapping that implements `__missing__`. In less than 20 lines of code, *autoconst.py* showcases the power of combining Python metaprogramming techniques

---

<sup>18</sup> The phrase is widely quoted. I found an early direct quote in [a post](#) in DHH’s blog from 2005.

I haven’t yet found a language that manages to be easy for beginners, practical for professionals, and exciting for hackers in the way that Python is. Thanks, Guido van Rossum and everybody else who makes it so.

## Further Reading

Caleb Hattingh—a technical reviewer of this book—wrote the *autoslot* package, providing a metaclass to automatically create a `__slots__` attribute in a user-defined class by inspecting the bytecode of `__init__` and finding all assignments to attributes of `self`. It’s useful and also an excellent example to study: only 74 lines of code in *autoslot.py*, including 20 lines of comments explaining the most difficult parts.

The essential references for this chapter in the Python documentation are “3.3.3. Customizing class creation” in the “Data Model” chapter of *The Python Language Reference*, which covers `__init_subclass__` and metaclasses. The *type class documentation* in the “Built-in Functions” page, and “4.13. Special Attributes” of the “Built-in Types” chapter in the *The Python Standard Library* are also essential reading.

In the *The Python Standard Library*, the *types module documentation* covers two functions added in Python 3.3 that simplify class metaprogramming: `types.new_class` and `types.prepare_class`.

Class decorators were formalized in *PEP 3129—Class Decorators*, written by Collin Winter, with the reference implementation authored by Jack Diederich. The PyCon 2009 talk “Class Decorators: Radically Simple” (*video*), also by Jack Diederich, is a quick introduction to the feature. Besides `@dataclass`, an interesting—and much simpler—example of a class decorator in Python’s standard library is `func tools.total_ordering` that generates special methods for object comparison.

For metaclasses, the main reference in Python’s documentation is *PEP 3115—Metaclasses in Python 3000*, in which the `__prepare__` special method was introduced.

*Python in a Nutshell*, 3rd ed., by Alex Martelli, Anna Ravenscroft, and Steve Holden, is authoritative, but was written before *PEP 487—Simpler customization of class creation* came out. The main metaclass example in that book—*MetaBunch*—is still valid, because it can’t be written with simpler mechanisms. Brett Slatkin’s *Effective Python*, 2nd ed. (Addison-Wesley) has several up-to-date examples of class building techniques, including metaclasses.



To learn about the origins of class metaprogramming in Python, I recommend Guido van Rossum’s paper from 2003, “[Unifying types and classes in Python 2.2](#)”. The text applies to modern Python as well, as it covers what were then called the “new-style” class semantics—the default semantics in Python 3—including descriptors and metaclasses. One of the references cited by Guido is *Putting Metaclasses to Work: a New Dimension in Object-Oriented Programming*, by Ira R. Forman and Scott H. Danforth (Addison-Wesley), a book to which he gave five stars on *Amazon.com*, adding the following review:

**This book contributed to the design for metaclasses in Python 2.2**

Too bad this is out of print; I keep referring to it as the best tutorial I know for the difficult subject of cooperative multiple inheritance, supported by Python via the `super()` function.<sup>19</sup>

If you are keen on metaprogramming, you may wish Python had the ultimate metaprogramming feature: syntactic macros, as offered in the Lisp family of languages and—more recently—by Elixir and Rust. Syntactic macros are more powerful and less error prone than the primitive code substitution macros in the C language. They are special functions that rewrite source code using custom syntax into standard code before the compilation step, enabling developers to introduce new language constructs without changing the compiler. Like operator overloading, syntactic macros can be abused. But as long as the community understands and manages the downsides, they support powerful and user-friendly abstractions, like DSLs (Domain-Specific Languages). In September 2020, Python core developer Mark Shannon posted [PEP 638—Syntactic Macros](#), advocating just that. A year after it was initially published, PEP 638 was still in draft and there were no ongoing discussions about it. Clearly it’s not a top priority for the Python core developers. I would like to see PEP 638 further discussed and eventually approved. Syntactic macros would allow the Python community to experiment with controversial new features, such as the walrus operator ([PEP 572](#)), pattern matching ([PEP 634](#)), and alternative rules for evaluating type hints (PEPs [563](#) and [649](#)) before making permanent changes to the core language. Meanwhile, you can get a taste of syntactic macros with the [MacroPy](#) package.

---

<sup>19</sup> I bought a used copy and found it a very challenging read.

## Soapbox

I will start the last soapbox in the book with a long quote from Brian Harvey and Matthew Wright, two computer science professors from the University of California (Berkeley and Santa Barbara). In their book, *Simply Scheme: Introducing Computer Science* (MIT Press), Harvey and Wright wrote:

There are two schools of thought about teaching computer science. We might caricature the two views this way:

1. **The conservative view:** Computer programs have become too large and complex to encompass in a human mind. Therefore, the job of computer science education is to teach people how to discipline their work in such a way that 500 mediocre programmers can join together and produce a program that correctly meets its specification.
2. **The radical view:** Computer programs have become too large and complex to encompass in a human mind. Therefore, the job of computer science education is to teach people how to expand their minds so that the programs can fit, by learning to think in a vocabulary of larger, more powerful, more flexible ideas than the obvious ones. Each unit of programming thought must have a big payoff in the capabilities of the program.

—Brian Harvey and Matthew Wright, preface to *Simply Scheme*<sup>20</sup>

Harvey and Wright’s exaggerated descriptions are about teaching computer science, but they also apply to programming language design. By now, you should have guessed that I subscribe to the “radical” view, and I believe Python was designed in that spirit.

The property idea is a great step forward compared to the accessors-from-the-start approach practically demanded by Java and supported by Java IDEs generating getters/setters with a keyboard shortcut. The main advantage of properties is to let us start our programs simply exposing attributes as public—in the spirit of *KISS*—knowing a public attribute can become a property at any time without much pain. But the descriptor idea goes way beyond that, providing a framework for abstracting away repetitive accessor logic. That framework is so effective that essential Python constructs use it behind the scenes.

Another powerful idea is functions as first-class objects, paving the way to higher-order functions. Turns out the combination of descriptors and higher-order functions enable the unification of functions and methods. A function’s `__get__` produces

---

<sup>20</sup> See p. xvii. Full text available at [Berkeley.edu](http://Berkeley.edu).